

Nimish CHAUDHARI

AI Infrastructure Engineer

Email: nimish.mailbox@gmail.com

Phone: +33 (0)7 63 55 45 51

Location: Luxembourg

LinkedIn: [nimishchaudhari](#)

GitHub: [nimishchaudhari](#)

CERTIFICATIONS & TRAINING

Microsoft Azure

- Azure Fundamentals (AZ-900)
- Azure Data Scientist (DP-100)


AI Security

- LLM Red Teaming & Jailbreak Benchmarking (PwC)

KEY ATTRIBUTES

- Client Relationship Management
- Fast Learner
- Technical Troubleshooting
- Cross-functional Collaboration
- Stakeholder Communication (Non-Technical)
- Technical Documentation

LANGUAGES



INTERESTS

- Starcraft 2
- Running
- Self hosting

PROFESSIONAL SUMMARY

AI Infrastructure Engineer focused on sovereign AI deployment and production-ready systems. I bridge local AI and cloud infrastructure — deploying LLMs on-premise and in-cloud with vLLM and llama.cpp, with security benchmarking and automation that balance performance, cost, and data sovereignty.

TECHNICAL SKILLS

CLOUD & INFRASTRUCTURE

- Azure, AWS, GCP
- Docker, Containerization
- Terraform, Ansible (IaC)
- CI/CD, DevOps Pipelines
- Proxmox, VM Networking, Coolify

AI ENGINEERING

- Python, TypeScript, Bash
- LLM Deployment (vLLM, llama.cpp)
- RAG, Prompt Orchestration
- LoRA/PEFT Fine-tuning
- Model Serving & Benchmarking

AI SECURITY

- AI Security Benchmarking & Red Teaming
- LLM Jailbreak Testing & Evaluation
- Prompt Injection Analysis
- Secure Cloud Infrastructure

PROJECTS

- **AI Security Benchmarking Framework** — Jailbreak evaluation across 3 benchmarks (300 prompts/model) on Qwen, Gemma, Minimax, Kimi — TTFT, throughput, and safety scores for sovereign AI provider selection
- **Self-Hosted LLM/HPC Infrastructure** — Containerized LLM services (vLLM, llama.cpp) with GPU optimization, quantization, and autoscaling
- **Proxmox Home Lab** — Self-hosted stack (Proxmox, Coolify, Docker): containerized services, VM networking, IaC workflows
- **Local AI Deployment Pipeline** — End-to-end sovereign AI: model selection, LoRA fine-tuning, inference optimization, monitoring

AI & DATA EXPERIENCE

Technology Consultant | AI & Cloud Infrastructure

PwC Advisory - Luxembourg

February 2023 – Present
3 years

- Led client-facing enterprise AI evaluations — RFI/RFP assessment, maturity analysis, vendor benchmarking; translated findings for non-technical stakeholders
- Deployed self-hosted LLM infrastructure (vLLM, llama.cpp) with GPU optimization, quantization, and containerized orchestration
- Built an AI security benchmarking framework for sovereign AI provider selection — jailbreak evaluation, TTFT/throughput, and safety scoring across open-weight models
- Designed LLM pipelines (RAG, serving, benchmarking) and automated infrastructure with AI-driven agents (Ansible, Docker, Coolify)

AI Development Intern | Computer Vision

Vodafone Procure & Connect (D2C Lab) - Luxembourg

October 2020 – September 2021
1 year

- Deep learning computer vision for electronic component identification — model design, AWS EC2 deployment, and a GUI for non-technical users
- Achieved **99.44% precision** in object detection for chipsets and motherboards
- Featured in European Business Review ([link](#))

Application Development Associate | Data Migration Specialist

Accenture France - Paris

April 2022 – September 2022
6 months

- Designed and modelled business processes using BPMN and automated workflows using Python/VBA
- Contributed to big data retail system implementation with large-scale legacy data migration

EDUCATION & RESEARCH

Master's in Information Systems (MIAGE)

Université Toulouse 1 Capitole, France

2019 – 2022

Coursework: AI Systems Architecture, Advanced Data Analysis, Business Intelligence

Bachelor's in Computer Science

University of Mumbai, India

2016 – 2019